

Class 7: Machine Learning

Mia Fava (A17577873)

Today we are going to learn how to apply different machine learning methods, beginning with clustering.

The goal here is to find groups/clusters in your input data.


First I will make up some data with clear groups. For this I will use the `rnorm()` function.

```
rnorm(10)
```

```
[1] -0.2676802  0.9789606  1.9138185  0.7371866  0.5999336  0.9330482  
[7]  1.0400379 -0.3929570 -2.3433195  0.5782594
```

```
hist(c(rnorm(10000, mean = -3), rnorm(1000, mean=3)))
```

Histogram of `c(rnorm(10000, mean = -3), rnorm(1000, mean = 3))`




```
n <- 30
x <- c(rnorm(n, -3), rnorm(n, +3))
y <- rev(x)
y
```

```
[1] 3.7356587 3.5873037 2.9045806 2.9661375 2.8697640 4.2568642
[7] 2.1334817 3.2846679 2.2451662 1.9947550 3.2701151 2.7267248
[13] 3.8053719 2.4114859 0.3464704 2.2869272 3.1224123 1.6861259
[19] 4.2069363 1.9110946 3.1471211 5.2920667 2.6190060 3.3643625
[25] 1.7771430 4.5619053 3.0200020 3.7186749 2.0649012 3.3955716
[31] -1.7272029 -0.8946366 -3.8617240 -3.9448530 -2.1146269 -4.5886410
[37] -2.5104449 -1.3819212 -3.5153242 -0.5345297 -4.0846879 -0.8310902
[43] -4.6958851 -3.7511674 -4.2599577 -1.8395624 -3.9798800 -2.8874919
[49] -2.4371982 -3.3690522 -4.4263811 -2.8112717 -3.8449749 -4.3716569
[55] -3.4405579 -3.9139590 -2.8742847 -2.8074599 -1.1206745 -4.2696506
```

```
z <- cbind(x, y)
head(z)
```

```
      x      y
[1,] -4.269651 3.735659
[2,] -1.120675 3.587304
[3,] -2.807460 2.904581
[4,] -2.874285 2.966138
[5,] -3.913959 2.869764
[6,] -3.440558 4.256864
```

```
plot(z)
```



Use `kmeans()` function setting `k` to 2 and `nstart = 20`. >Q How many points are in each cluster

Q What ‘component’ of your result object details? -cluster size -cluster assignment/member -cluster center?

Q. Plot x colored by the kmean cluster assignment and add cluster centers as blue points

```
km <- kmeans(z, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.036358	2.957093
2	2.957093	-3.036358

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 73.69228 73.69228
      (between_SS / total_SS =  88.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Reults in kmeans object km

```
attributes(km)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
$class
[1] "kmeans"
```

Cluster size?

km\$size

[1] 30 30

cluster assignment/member?

```
km$cluster
```

[illegible]

- cluster center?


km\$centers

	x	y
1	-3.036358	2.957093
2	2.957093	-3.036358


Q Plot z colored by the kmeans cluster assignment and add cluster centers as blue points

R will recycle the shorter color version to be the same length as the longer

```
plot(z, col=c("red", "blue"))
```




```
plot(z, col=km$cluster)
points(km$cluster, col="blue", pch=17, cex=3)
```



Q. Can you run kmean and ask for 4 clusters, please and plot the results like we have done above?

```
pm <- kmeans(z, centers=4)
```

```
plot(z, col=km$cluster)
points(km$cluster, col="blue", pch=15, cex=1)
```



##Hierarchial Clustering

Lets take the same made-up datat z and see how hclust works


```
d<- dist(z)
hc <-hclust(d)
hc
```

Call:
hclust(d = d)

Cluster method : complete
Distance : euclidean
Number of objects: 60

```
plot(hc)
abline(h=8, col="red")
```

Cluster Dendrogram



```
cutree(hc, h=8)
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 3 3 3 3 3 3 3 3
[39] 3 3 3 3 3 3 3 2 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```


I can get my membership c=vector by “cutting the tree” with the `cutree()` function like so:

```
grps<-cutree(hc, h=8)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 3 3 3 3 3 3 3 3
[39] 3 3 3 3 3 3 3 2 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
```

Can you plot z colored by our hclust results:

```
plot(z, col=grps)
```



##START OF ON HANDS LAB

PCA of UK food data

Read the data from the UK on food consumption in different parts of the UK

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17  4
```

```
##Preview the first 6 rows
```

```
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
# Note how the minus indexing works
```

```
rownames(x) <- x[,1]
```

```
x <- x[,-1]
```

```
head(x)
```

	Wales	Scotland	N.Ireland
105	103	103	66
245	227	242	267
685	803	750	586
147	160	122	93
193	235	184	209
156	175	147	139

```
dim(x)
```

```
[1] 17  3
```

```
x <- read.csv(url, row.names=1)
```

```
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139


Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

Personally I prefer the `read.csv()` function to be the easiest way to get the row names in comparison to the `view()` function.


Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

Changing the `beside =` from T to F would change the first plot to the second plot.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

This following code for the pairs plot is useful with small datas like the one provided. It is also showing the country names in a diagonal – in which the row is showing what country is on the y-axis; the x axis shows whatever is being compared. If the data is the same for the two countries, the data will look like a straight diagonal line; if they are not similar/differ in data value, they will be more scattered and not as uniform across a diagonal line. It is hard to see structure and trend in this plotting, however we have to do this when we have big datasets with 1,000s or 10s of thousands of things we are measuring.

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

The main differences between N. Ireland and other countries of the UK is that Irelands data values are seen to be more scattered, therefore the data values in each category of food are seen to be less similar to those of the other UK countries.

PCA to the rescue

See how PCA deals with this data set – main function in base R to do PCA is called `prcomp()`

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is inside this `pca` object that we created from running `prcomp()`

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```


```
$class
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13


Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1(67.%)", ylab="PC2(29%)")
text(pca$x[,1], pca$x[,2], colnames(x), col=c("blue", "red", "orange", "green"))
```




```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```


```
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	2.921348e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```




```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

Changed the `pca$rotation[,1]` to `topcarotation[,2]`. The two groups that feature prominently were the high negative loading score of the fresh potatoes in Scotland predominantly, and soft drinks are the positive loading score seen in N.Ireland, Wales, & England.

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```



Q10 :